

# Database Systems

---

## *Course Project Instruction*

Ma Dongzhe / 马冬哲  
152 1062 0224  
mdzfirst@gmail.com

# Policy

- 2-3 persons form a team.
- 60% of your final score.
- In December, 3-4 outstanding teams will be invited to make a presentation.

# **The Task is**

To Implement a DBMS Prototype.

# What We Care

- Correctness
- Response Time
  - Storage
  - Access Method
  - Caching Strategy
  - Optimizing
  - Query Processing

# **What We Don't Care**


- Transaction Processing
- Concurrency Control
- Crash Recovery

# Grading Criteria

Accomplishment	At least one correct run	10
Overall Evaluation	Correctness & Design & Code Quality & Contrib.	10
Performance	$S_j = \frac{\text{sum}((T_{i,\text{best}} / T_{i,j})^{0.2})}{\text{Full}} * (S_j / S_{\text{best}})$	30
Documentation	Content & Feature	10
Presentation	For some teams only	$\leq 5$

# Example

	Workload 0	Workload 1
Team 0	5	100
Team 1	10	1000
Team 2	1	Fail



$$S_0 = (1 / 5)^{0.2} + (100 / 100)^{0.2} = 1.725$$

$$S_1 = (1 / 10)^{0.2} + (100 / 1000)^{0.2} = 1.262$$

$$S_2 = (1 / 1)^{0.2} + (100 / \text{INF})^{0.2} = 1$$

$$\text{Score}_0 = 30 * 1.725 / 1.725 = 30$$

$$\text{Score}_1 = 30 * 1.262 / 1.725 = 22$$

$$\text{Score}_2 = 30 * 1 / 1.725 = 17$$

$$(1 / 5)^{0.2} = 0.725$$

$$(1 / 10)^{0.2} = 0.631$$

$$(1 / 50)^{0.2} = 0.457$$

$$(1 / 100)^{0.2} = 0.398$$

$$(1 / 500)^{0.2} = 0.289$$

$$(1 / 1000)^{0.2} = 0.251$$

$$(1 / 5000)^{0.2} = 0.182$$

# The Environment is

- Ubuntu 10.04 LTS, 32-bit
- g++ 4.4.3
- Intel(R) Xeon(R) 5130 @ 2.00GHz x2
- 3.0 GB RAM, 1.9 GB swap



Nonnegotiable



# You Have to Implement

- **create()**

*Create a new table.*

- **train()**

*Given some query information, train your system and choose the storage and access methods.*

- **load()**

*Load initial data in csv format. The initial data set might be too large to keep in the main memory entirely.*

- **preprocess()**

*Build the indexes and do other preprocessing.*

*See [course/include/client.h](#) for more details.*

# You Have to Implement

- `execute()`  
*Execute a query or insert statement.*
- `next()`  
*Get the next row from the result set of the last query.*
- `close()`  
*Close the sockets and kill other threads.*

*See `course/include/client.h` for more details.*

# You Have to Implement

- **execute()**

*Execute a query or insert statement.*

- **next()**

*Get the next row from the result set of the last query.*


- **close()**

*Close the sockets and kill other threads.*

**WARNING: Run time of execute()  
and next() will be measured.**

*See [course/include/client.h](#) for more details.*

# Test Procedure



# Query Statement

```
SELECT column0, column1, ...  
FROM table0, table1, ...  
WHERE condition0 AND ... AND conditionN;
```

A condition could be

```
column = constant  
column < constant (For integers only)  
column > constant (For integers only)  
column0 = column1 (Join condition)
```

# Query Statement

```
SELECT column0, column1, ...  
FROM table0, table1, ...  
WHERE condition0 AND ... AND conditionN;
```

A condition could be

No prefix

column = constant

column < constant (For integers only)

column > constant (For integers only)

column0 = column1 (Join condition)

# Query Statement

```
SELECT column0, column1, ...  
FROM table0, table1, ...  
WHERE condition0 AND ... AND conditionN;
```

A condition could be

column = constant

column < constant (For integers only)

column > constant (For integers only)

column0 = column1 (Join condition)

} Same type

# Query Statement

```
SELECT column0, column1, ...  
FROM table0, table1, ...  
WHERE condition0 AND ... AND conditionN;
```

A condition could be

The only operator



column = constant

column < constant (For integers only)

column > constant (For integers only)

column0 = column1 (Join condition)



# Query Statement

```
SELECT column0, column1, ...
```

```
FROM table0, table1, ...
```

```
WHERE condition0 AND ... AND conditionN;
```

A condition could be

If the FROM-clause contains only one table, there might be no WHERE-clause.

```
column = constant
```

```
column < constant (For integers only)
```

```
column > constant (For integers only)
```

```
column0 = column1 (Join condition)
```

# Insert Statement

```
INSERT INTO table  
VALUES (value_list0), ..., (value_listN);
```

All value lists are in csv format.

```
constant0,constant1,...,constantN
```

# Insert Statement

```
INSERT INTO table  
VALUES (value_list0), ..., (value_listN);
```



All value lists are in csv format.

No column list

```
constant0, constant1, ..., constantN
```

# Insert Statement

```
INSERT INTO table  
VALUES (value list0), ..., (value listN);
```

All value lists are in csv format. **Number of rows is important for the train() routine.**

```
constant0, constant1, ..., constantN
```

# Format

```
SELECT _a_, _b_  
FROM _A_, _B_  
WHERE _a_ = _5_ AND _b_ < _10_;
```

```
INSERT _INTO _A_ VALUES_  
(_0, 'Stalin', 1879_)_,_  
(_1, 'Roosevelt', 1882_)_;
```

# Data Types

- **INTEGER**

*32-bit unsigned integer, 'int' is OK.*

- **VARCHAR(d)**

*Consist of \_, a-z, A-Z, or 0-9. Enclosed by single quotes.  
At most d characters (excluding the quotes).*

## **NOTE:**

*All identifiers (table names and column names) are string constants not starting with 0-9.*

*Columns in different tables have distinct names.*

*String constants don't contain space, quote, or comma.*

# Primary Keys

- Primary keys will be assigned to all relations.
- The primary keys will be unique. There is no need to check this constraint.
- The primary keys will be given in ascending order.
- You can just ignore them.

# Join Operations

Let nodes represent tables and edges represent join conditions, then each query can be transformed into a graph. This graph should be a tree which

- is connected;
- contains no self-cycles;
- contains no duplicate edges;
- contains no cycles (at least 3 nodes).



# Workloads

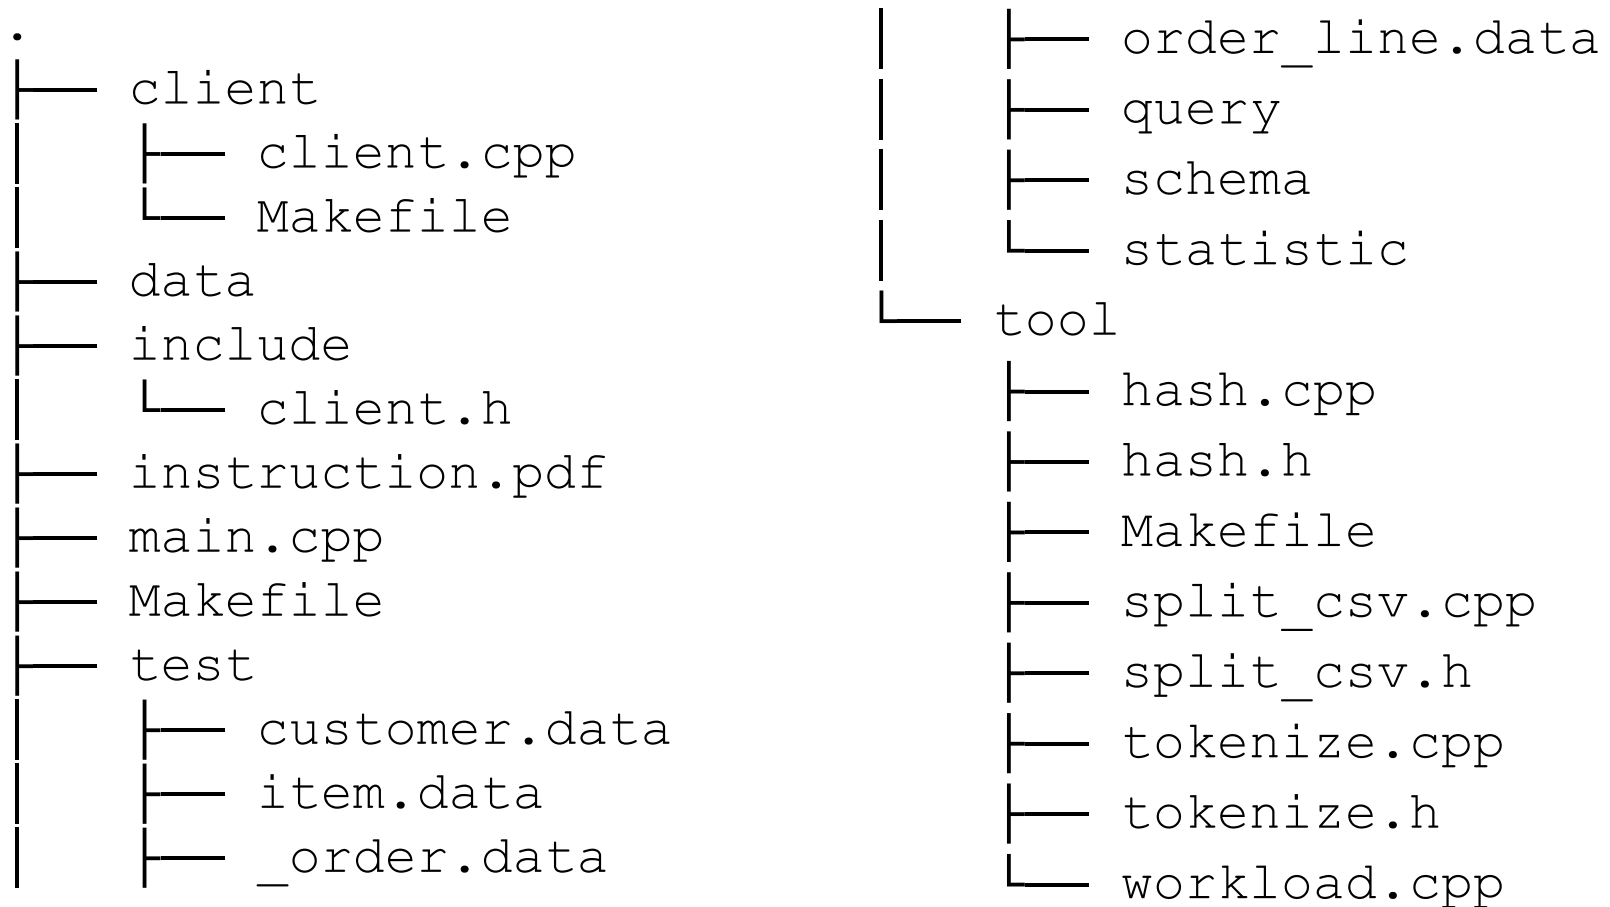
- Projection
  - Selection
  - Join
- 
- TPC-C
  - TPC-H

# Workloads

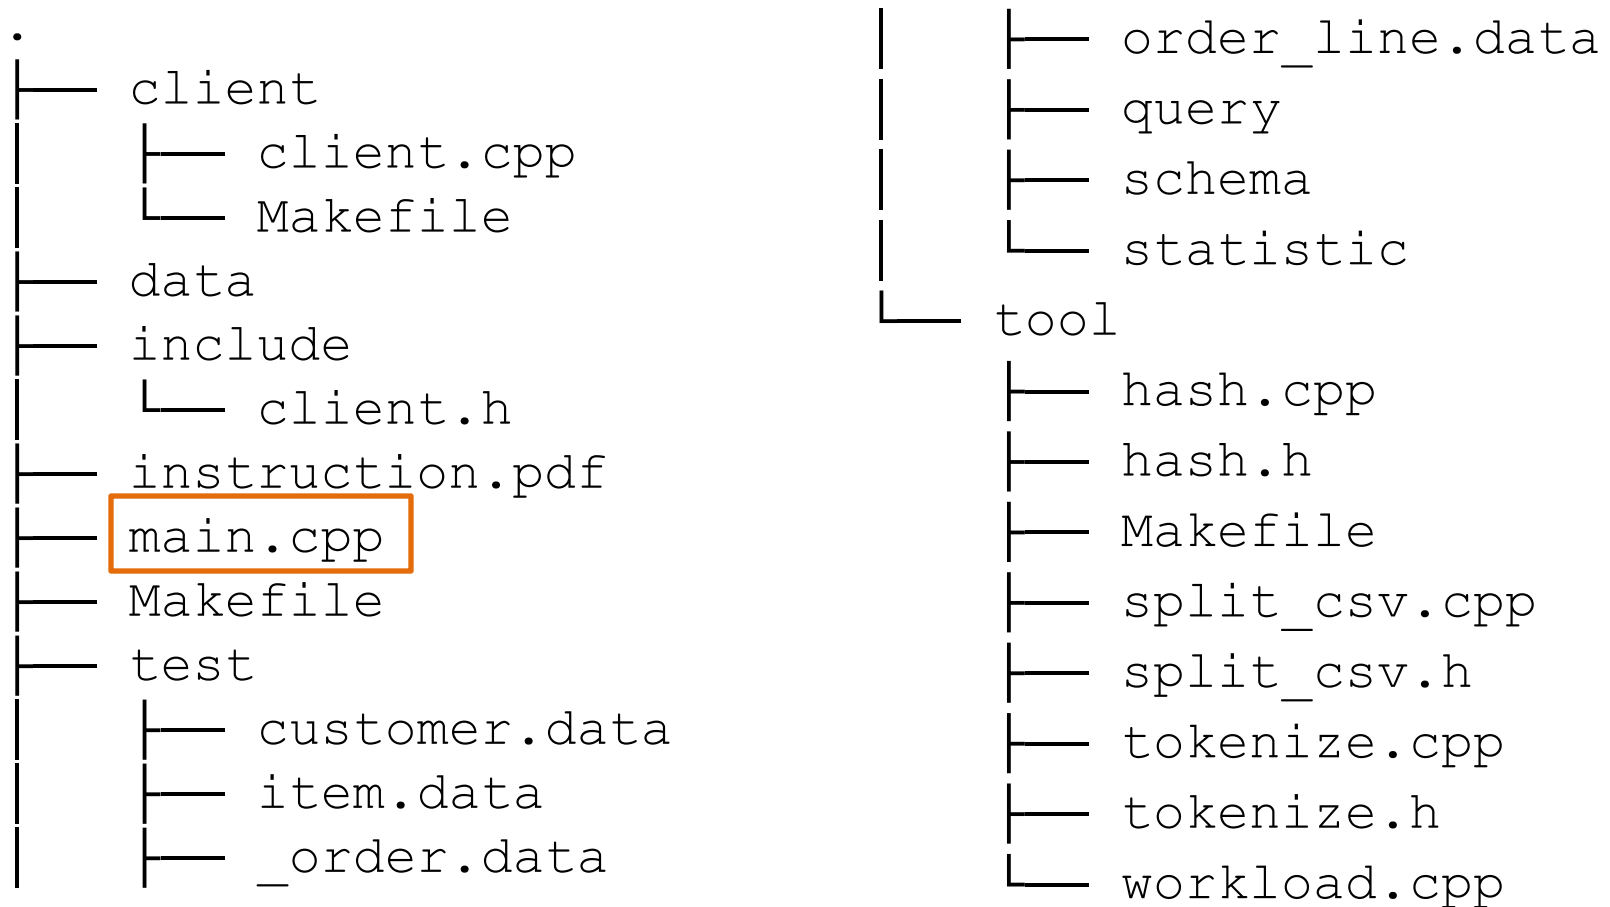
- Projection
  - Selection
  - Join
- 
- TPC-C
  - TPC-H

```
string workload();
```

# Directory Structure

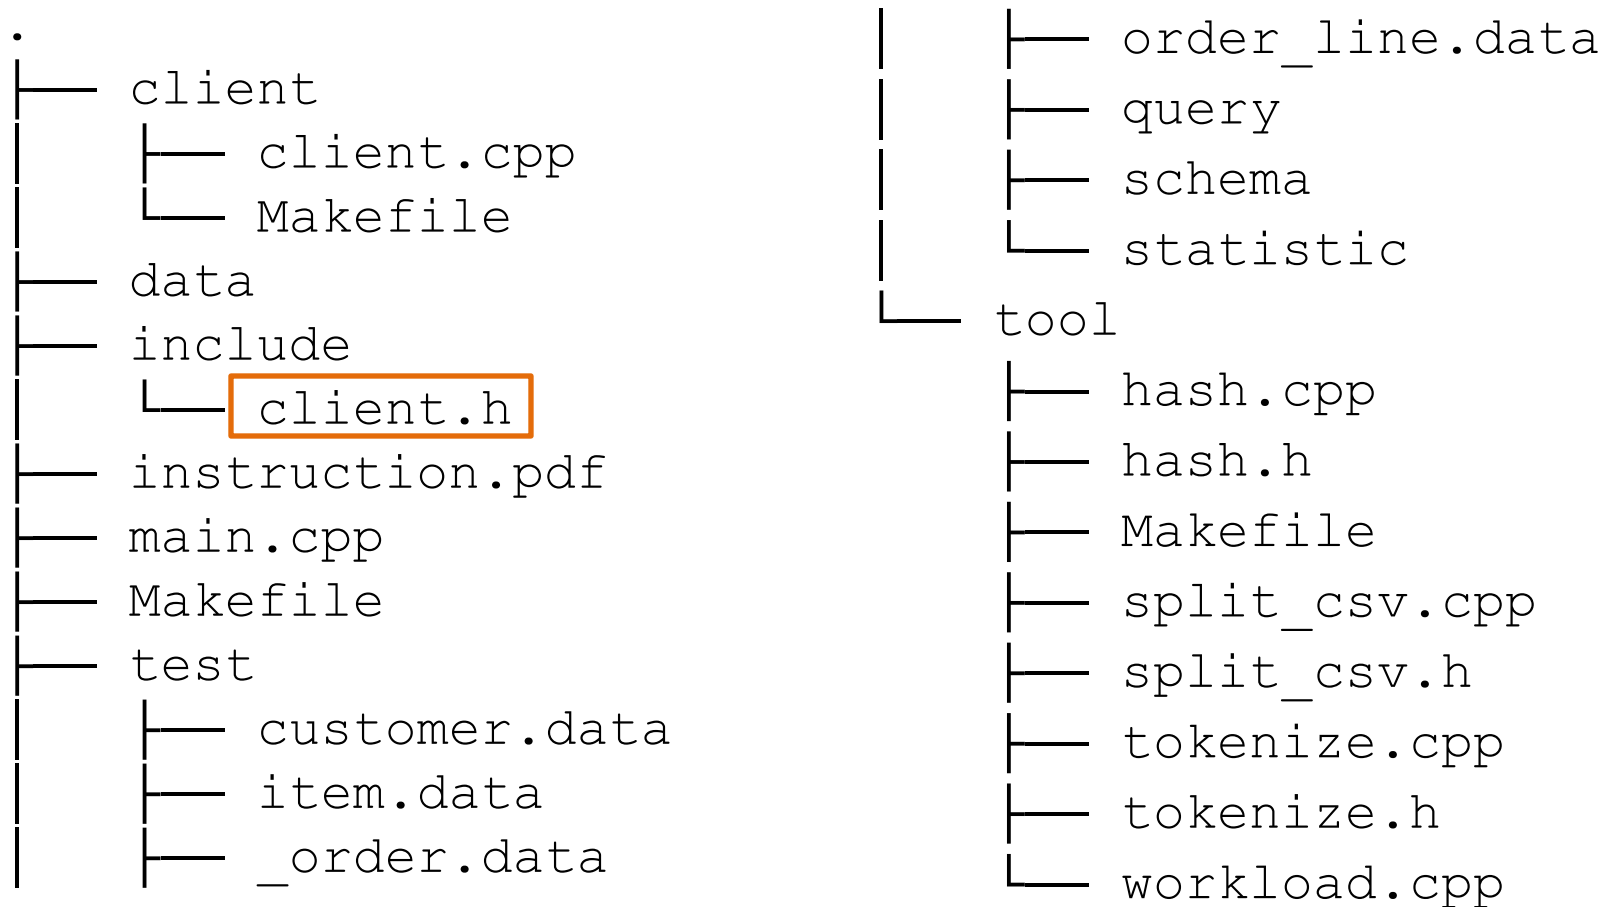


# Directory Structure



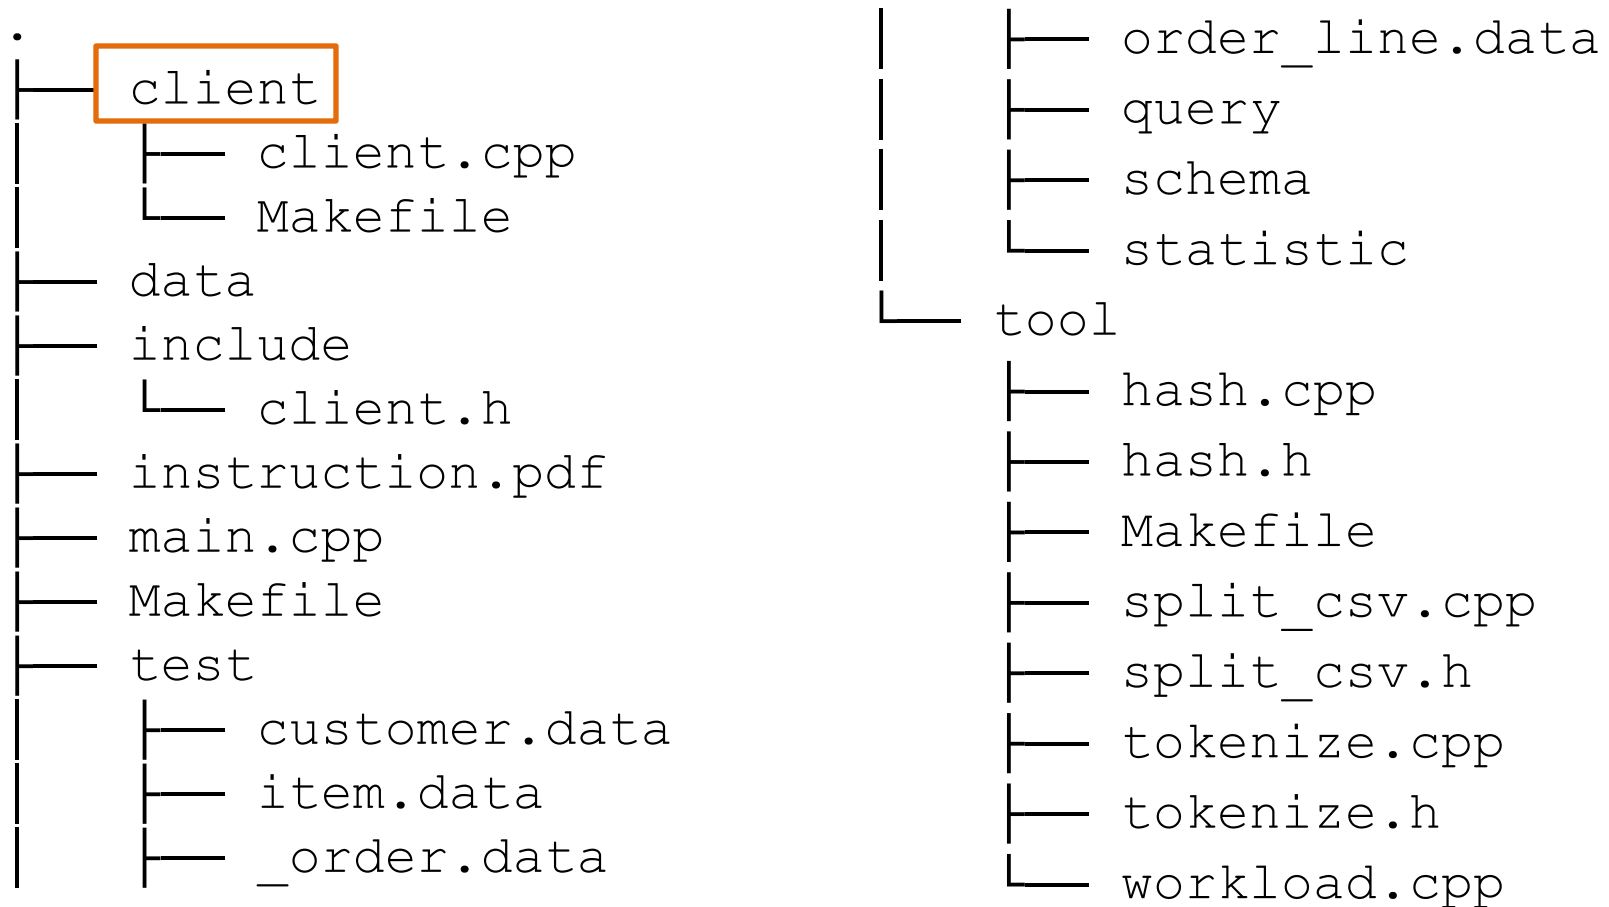
*Test procedure. This file will be modified.*

# Directory Structure



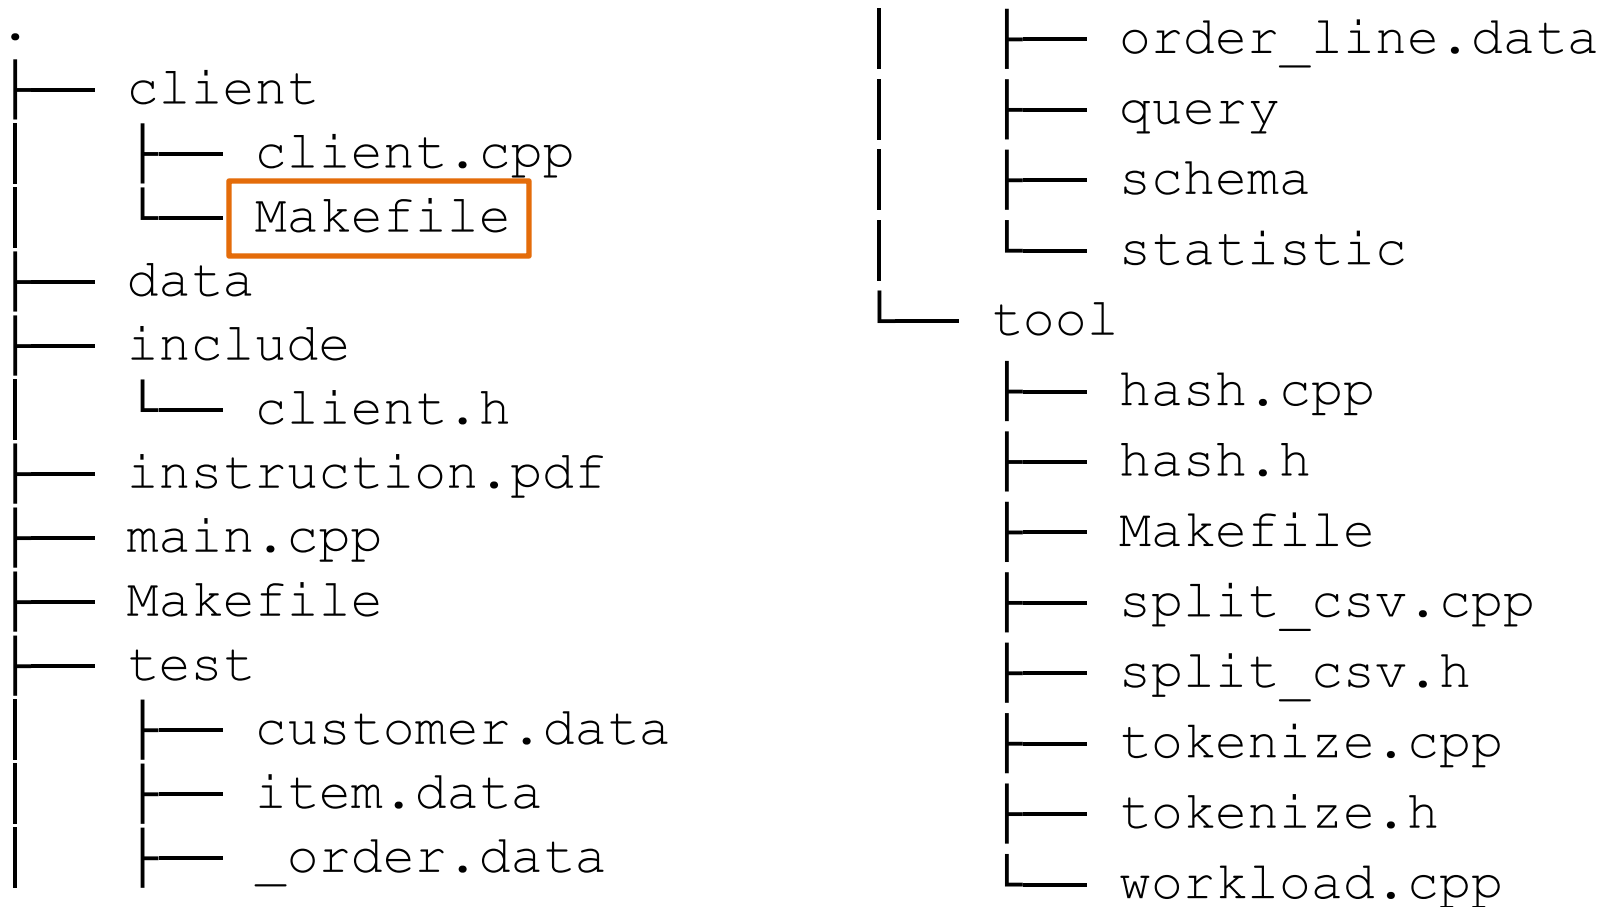
*APIs to implement.*

# Directory Structure



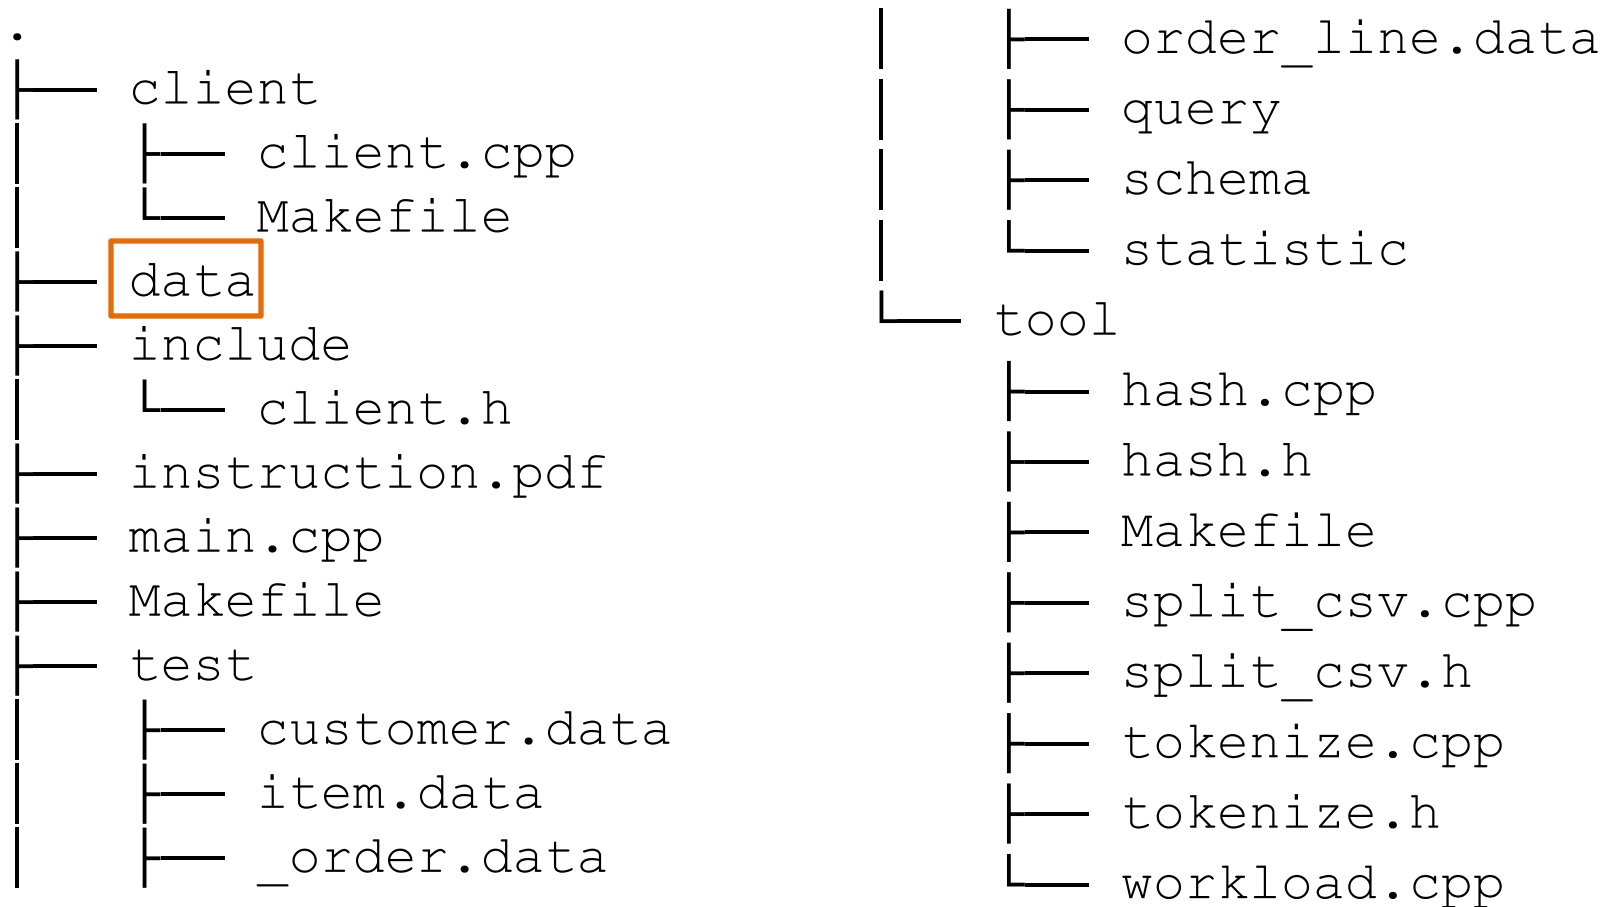
*Keep all your source codes in this directory.*

# Directory Structure



*Make sure your Makefile is correct.*

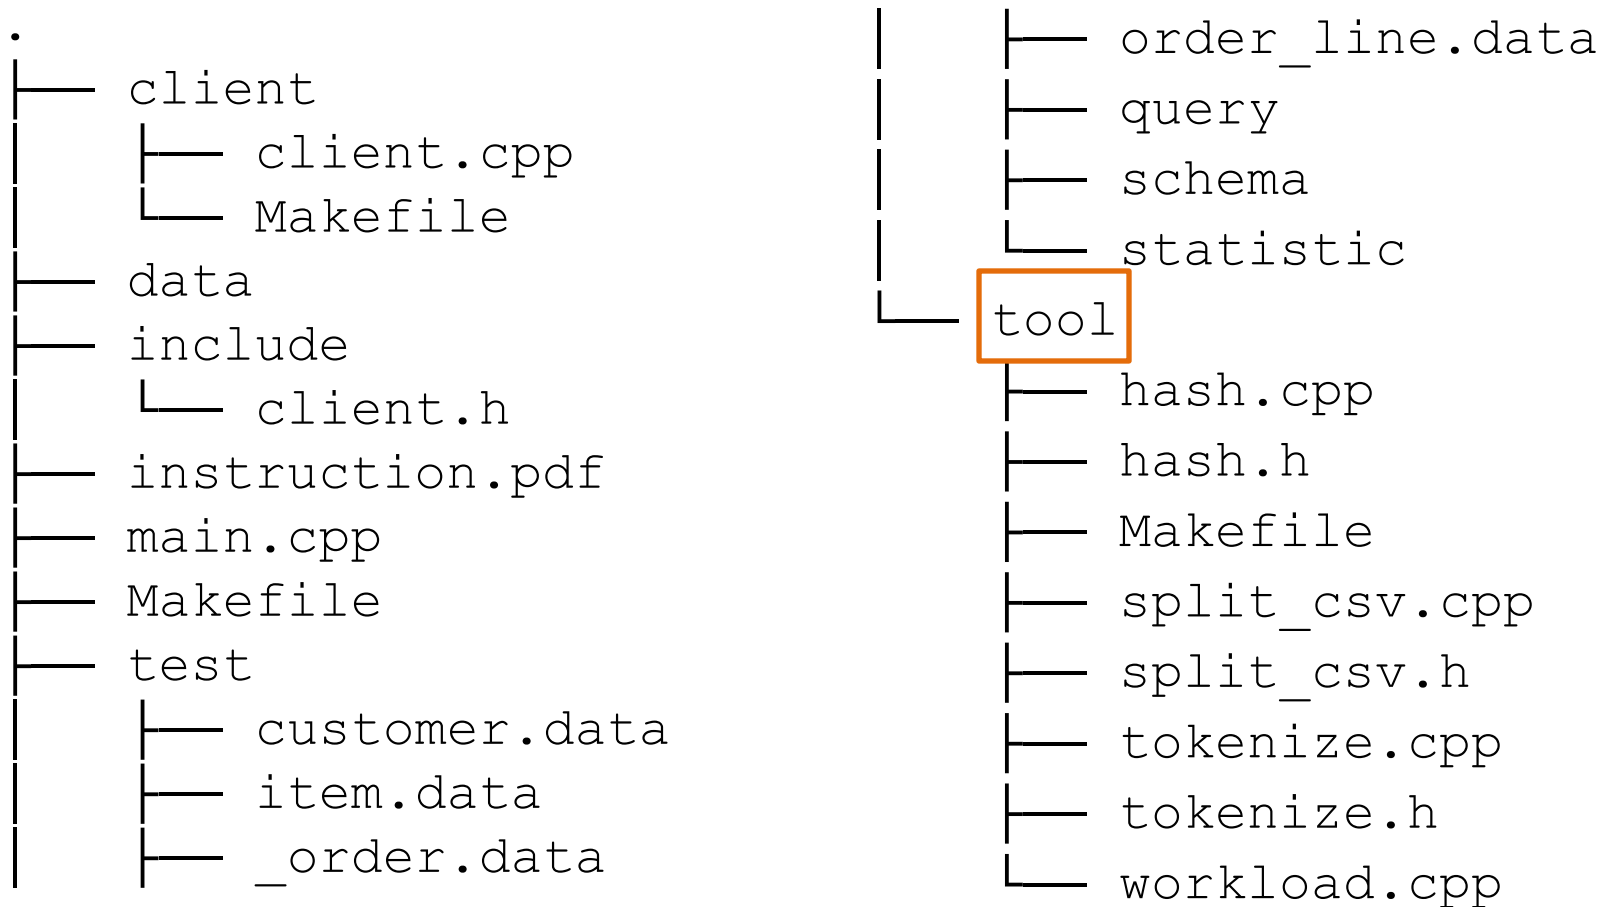
# Directory Structure



*Keep all your data in this directory.*



# Directory Structure



*Do NOT use any routines in other folders.*

# About Third-party Library

- You are free to use any third-party library or code about storage, index, multi-thread, network, etc.

*e.g. Boost, Berkeley DB, open-source disk-based B-tree / hash table implementation, etc*

- You are forbidden to use any system that is capable to process a SQL query.

*e.g. MySQL, PostgreSQL, etc*

- Ask for confirmation if you are not sure.

# Document / Presentation

- System Architecture
- Storage Model and the Selection Strategy
- Index Structure and the Selection Strategy
- Caching Strategy
- Query Processing Strategy
  - Heuristic Rules
  - Cost Model
- Other features of your system
- References
- Personal Contribution Rate (For document)

# Submission

- Prepare your submission with *make tar*.
- Submit the *\*.tar.gz* to ...
- One submission every 2 hours.
- Only the last submission counts.

# Notice

- Some tests may depend on the result of another one.
- Time limitation applies to the whole program, although only `execute()` and `next()` affect your final scores.
- Not all the queries are provided in `train()`, although only those included are measured.

# Warnings

- Never do irrelevant operations
- Never replicate other team's work



# Hints

- Read some books and research papers
- Discuss with others
- Start **ASAP**

# **create()**

Keep the schema safe.



# **train()**

- Find affinitive tables.
- Find affinitive attributes.
- Choose access methods.
- Read-intensive or update-intensive?

# **load()**

Keep the data safe.

# **preprocess()**

- Make some useful statistics.
- Build some indexes.
- Start some threads.

# Statistics

- For uniform distribution
  - $\text{Size}(R)$ ,  $\text{Cnt}(R)$ ,  $\text{Card}(A)$ ,  $\text{Min}(A)$ ,  $\text{Max}(A)$
  - $\text{SF}(A = \text{value}) = 1 / \text{Card}(A)$
  - $\text{SF}(A > \text{value}) = (\text{Max}(A) - \text{value}) / \text{Range}(A)$
  - $\text{SF}(A < \text{value}) = (\text{value} - \text{Min}(A)) / \text{Range}(A)$
  - $\text{SF}(A_0 \wedge A_1) = \text{SF}(A_0) * \text{SF}(A_1)$
- For skewed distribution (e.g., Zipf)
  - Histogram

# **execute() and next()**

- Do all the jobs in execute().
- Do all the jobs in next().
- Do all the jobs in independent threads.

# Query Processing

- Google 'query processing'
- Search on ACM Digital Library (*dl.acm.org*)
- Cost model vs. Rules




# Join Operation


- Nested Loop Join
- Index-based Nested Loop Join
- Sort-Merge Join
- Hash Join (Pruning)

# **close()**

Close your program safely.







大学计算机教育国外著名教材、教参系列 (影印版)

# Principles of Distributed Database Systems

PEARSON  
Prentice  
Hall

M. Tamer Özsu  
Patrick Valduriez

Second Edition

## 分布式数据库 系统原理

(第 2 版)



清华大学出版社

Good luck and have fun!